

Application of Information Theory, Lecture 7

Relative Entropy

Handout Mode

Iftach Haitner

Tel Aviv University.

December 1, 2015

Part I

Statistical Distance

Statistical distance

- ▶ Let $p = (p_1, \dots, p_m)$ and $q = (q_1, \dots, q_m)$ be distributions over $[m]$
- ▶ Their **statistical distance** (also known as, variation distance) is defined by

$$\text{SD}(p, q) := \frac{1}{2} \sum_{i \in [m]} |p_i - q_i|$$

- ▶ This is simply the L_1 norm between the distribution vectors
- ▶ We will see other “distance” measures for distributions next lecture
- ▶ For $Z \sim p$ and $Y \sim q$, let $\text{SD}(X, Y) = \text{SD}(p, q)$
- ▶ Claim (HW): $\text{SD}(p, q) = \max_{S \subseteq [m]} (\sum_{i \in S} p_i - \sum_{i \in S} q_i)$
- ▶ Hence, $\text{SD}(p, q) = \max_D (\Pr_{X \sim p} [D(X) = 1] - \Pr_{X \sim q} [D(X) = 1])$
- ▶ Interpretation

Distance from the uniform distribution

- ▶ Let X be rv over $[m]$
- ▶ $H(X) \leq \log m$
- ▶ $H(X) = \log m \iff X$ is uniform over $[m]$

Theorem 1 (this lecture)

Let X rv over $[m]$. Assume $H(X) \geq \log m - \varepsilon$, then

$$\text{SD}(X, \sim [m]) \leq \sqrt{\varepsilon \cdot \frac{\ln 2}{2}} = O(\sqrt{\varepsilon})$$

Part II

Relative entropy Distance

Section 1

Definition and Basic Facts

Definition

- ▶ For $p = (p_1, \dots, p_m)$ and $q = (q_1, \dots, q_m)$, let

$$D(p||q) = \sum_{i=1}^m p_i \log \frac{p_i}{q_i}$$

$$0 \log \frac{0}{0} = 0, p \log \frac{p}{0} = \infty$$

- ▶ The relative entropy of pair of rv's, is the relative entropy of their distributions.
- ▶ Names: Entropy of p relative to q , relative entropy, information divergence, Kullback-Leibler (KL) divergence/distance
- ▶ Many different interpretations
- ▶ Main interpretation: the information we **gained** about X , if we originally thought $X \sim q$ and now we learned $X \sim p$

Numerical Example

$$D(p\|q) = \sum_{i=1}^m p_i \log \frac{p_i}{q_i}$$

► $p = (\frac{1}{4}, \frac{1}{2}, \frac{1}{4}, 0), q = (\frac{1}{2}, \frac{1}{4}, \frac{1}{8}, \frac{1}{8})$

► $D(p\|q) = \frac{1}{4} \log \frac{\frac{1}{4}}{\frac{1}{2}} + \frac{1}{2} \log \frac{\frac{1}{2}}{\frac{1}{4}} + \frac{1}{4} \log \frac{\frac{1}{4}}{\frac{1}{8}} + 0 \log 0 = \frac{1}{4} \cdot (-1) + \frac{1}{2} \cdot 1 + \frac{1}{4} \cdot 1 = \frac{1}{2}$

► $D(q\|p) = \frac{1}{2} \log \frac{\frac{1}{2}}{\frac{1}{4}} + \frac{1}{4} \log \frac{\frac{1}{4}}{\frac{1}{2}} + \frac{1}{8} \log \frac{\frac{1}{8}}{\frac{1}{4}} + \frac{1}{8} \log \frac{\frac{1}{8}}{0} =$
 $\frac{1}{2} \cdot 1 + \frac{1}{4} \cdot (-1) + \frac{1}{8} \cdot (-1) + \infty = \infty$

Supporting the interpretation

- ▶ X rv over $[m]$
- ▶ $H(X)$ — measure for amount of information we do not have about X
- ▶ $\log m - H(X)$ — measure for information we **do** have about X (just by knowing its distribution)
- ▶ Example $X = (X_1, X_2) \sim (\frac{1}{2}, 0, 0, \frac{1}{2})$ over $\{00, 01, 10, 11\}$
- ▶ $H(X) = 1$, $\log m - H(X) = 2 - 1 = 1$
- ▶ Indeed, we know $X_1 \oplus X_2$

▶

$$\begin{aligned} H(\sim [m]) - H(p_1, \dots, p_m) &= \log m - H(p_1, \dots, p_m) \\ &= \log m + \sum_i p_i \log p_i = \sum_i p_i (\log p_i - \log \frac{1}{m}) \\ &= \sum_i p_i \log \frac{p_i}{\frac{1}{m}} = D(p \| \sim [m]) \end{aligned}$$

- ▶ $D(X \| \sim [m])$ — **measures** the information we **gained** about X , if we originally thought it is $\sim [m]$ and now we learned it is $\sim p$

Supporting the interpretation, cont.

- ▶ (generally) $D(p\|q) \neq H(p) - H(q)$
- ▶ $H(p) - H(q)$ is **not** a good measure for information change
- ▶ Example: $q = (0.01, 0.99)$ and $p = (0.99, 0.01)$
- ▶ We were almost sure that $X = 1$ but learned that X is almost surely 0
- ▶ But $H(p) - H(q) = 0$
- ▶ Also, $H(p) - H(q)$ might be negative
- ▶ We **understand** $D(p\|q)$ as the information we gained about X , if we originally thought it is $\sim q$ and now we learned it is $\sim p$

Changing distribution

- ▶ What does it mean: originally thought $X \sim q$ and now we learned $X \sim p$?

How can a distribution change?

- ▶ Typically, this happens by learning additional information
- ▶ $q_i = \Pr[X = i]$ and $p_i = \Pr[X = i|E]$
- ▶ Example $X \sim (\frac{1}{2}, \frac{1}{4}, \frac{1}{4}, 0)$; someone saw X and tells us that $X \leq 2$
- ▶ The distribution changes to $X \sim (\frac{2}{3}, \frac{1}{3}, 0, 0)$

- ▶ Another example

$X \backslash Y$	1	2	3	4
0	$\frac{1}{4}$	$\frac{1}{4}$	0	0
1	$\frac{1}{4}$	0	$\frac{1}{4}$	0

- ▶ $Y \sim (\frac{1}{2}, \frac{1}{4}, \frac{1}{4}, 0)$, but
- ▶ $Y \sim (\frac{1}{2}, \frac{1}{2}, 0, 0)$ conditioned on $X = 0$
- ▶ $Y \sim (\frac{1}{2}, 0, \frac{1}{2}, 0)$ conditioned on $X = 1$
- ▶ Generally, a distribution can change if we condition on event E

Additional properties

- ▶ $0 \log \frac{0}{0} = 0$, $p \log \frac{p}{0} = \infty$ for $p > 0$
- ▶ $\exists i$ s.t. $p_i > 0$ and $q_i = 0$, then $D(p\|q) = \infty$
- ▶ If originally $\Pr[X = i] = 0$, then it cannot be more than 0 after we learned something.
- ▶ Hence, it make sense to think of it as infinite amount of information learnt
- ▶ Alternatively, we can define $D(p\|q)$ only for distribution with $q_i = 0 \implies p_i = 0$
(recall that $\Pr[X = i] = 0 \implies \Pr[X = i|E] = 0$, for any event E)
- ▶ If p_i is large and q_i is small, then $D(p\|q)$ is large
- ▶ $D(p\|q) \geq 0$, with equality iff $p = q$ (hw)

Example

- ▶ $q = (q_1, \dots, q_m)$ with $\sum_{i=1}^n q_i = 2^{-k}$ (i.e., $n < m$)
- ▶ $p_i = \begin{cases} q_i/2^{-k}, & 1 \leq i \leq n \\ 0, & \text{otherwise.} \end{cases}$
- ▶ $p = (p_1, \dots, p_m)$ — the distribution of q conditioned on the event $i \in [n]$
- ▶ $D(p||q) = \sum_{i=1}^n p_i \log \frac{p_i}{q_i} = \sum_{i=1}^n p_i \log 2^k = \sum_{i=1}^n p_i k = k$
- ▶ We gained k bits of information
- ▶ Example: $\sum_{i=1}^n q_i = \frac{1}{2}$, and we were told that $i \leq n$ or $i > n$, we got one bit of information

Section 2

Axiomatic Derivation

Axiomatic derivation

Let \tilde{D} is a continuous and symmetric (wrt each distribution) function such that

1. $\tilde{D}(p \parallel \sim [m]) = \log m - H(p)$
2. $\tilde{D}((p_1, \dots, p_m) \parallel (q_1, \dots, q_m)) = \tilde{D}((p_1, \dots, p_{m-1}, \alpha p_m, (1 - \alpha)p_m) \parallel (q_1, \dots, q_{m-1}, \alpha q_m, (1 - \alpha)q_m))$, for any $\alpha \in [0, 1]$

then $\tilde{D} = D$.

Interpretation

Proof: Let p and q be distributions over $[m]$, and assume $q_i \in \mathbb{Q} \setminus \{0\}$.

► $\tilde{D}(p \parallel q) = \tilde{D}((\alpha_{1,1}p_1, \dots, \alpha_{1,k_1}p_1, \dots, \alpha_{m,1}p_m, \dots, \alpha_{m,k_m}p_m) \parallel (\alpha_{1,1}q_1, \dots, \alpha_{1,k_1}q_1, \dots, \alpha_{m,1}q_m, \dots, \alpha_{m,k_m}q_m))$, for $\sum_j \alpha_{i,j} = 1$ and $\alpha_{i,j} \geq 0$

► Taking α 's s.t. $\alpha_{i,1} = \alpha_{i,2} \dots, \alpha_{i,k_i} = \alpha_i$ and $\alpha_i q_i = \frac{1}{M}$, it follows that

$$\begin{aligned}\tilde{D}(p \parallel q) &= \log M - H((\alpha_{1,1}p_1, \dots, \alpha_{1,k_1}p_1, \dots, \alpha_{m,1}p_m, \dots, \alpha_{m,k_m}p_m)) \\ &= \sum p_i \log M + \sum_i p_i \log \alpha_i p_i = \sum_i p_i (\log M + \log \frac{p_i}{q_i M}) = \sum_i p_i \log \frac{p_i}{q_i}.\end{aligned}$$

► Zeros and non-rational q_i 's are dealt by continuity

Section 3

Relation to Mutual Information

Mutual information as expected relative entropy

Claim 2

$$\mathbb{E}_{y \leftarrow Y} [D(X|_{Y=y} \| X)] = I(X; Y).$$

Proof:

- ▶ Let $X \sim (q_1, \dots, q_m)$ over $[m]$, and Y be rv over $\{0, 1\}$
- ▶ $(X|_{Y=j}) \sim p_j = (p_{j,1}, \dots, p_{j,m})$, $p_{j,i} = \Pr[X = i | Y = j]$

$$\begin{aligned} \mathbb{E}_Y [D(p_Y \| q)] &= \Pr[Y = 0] \cdot D(p_{0,1}, \dots, p_{0,m} \| q_1, \dots, q_m) \\ &\quad + \Pr[Y = 1] \cdot D(p_{1,1}, \dots, p_{1,m} \| q_1, \dots, q_m) \\ &= \Pr[Y = 0] \cdot \sum_i p_{0,i} \log \frac{p_{0,i}}{q_i} + \Pr[Y = 1] \cdot \sum_i p_{1,i} \log \frac{p_{1,i}}{q_i} \\ &= \Pr[Y = 0] \cdot \sum_i p_{0,i} \log p_{0,i} + \Pr[Y = 1] \cdot \sum_i p_{1,i} \log p_{1,i} \\ &\quad - \Pr[Y = 0] \cdot \sum_i p_{0,i} \log q_i - \Pr[Y = 1] \cdot \sum_i p_{1,i} \log q_i \\ &= -H(X|Y) - \sum_i (\Pr[Y = 0] \cdot p_{0,i} + \Pr[Y = 1] \cdot p_{1,i}) \log q_i \\ &= -H(X|Y) + H(X) = I(X; Y). \square \end{aligned}$$

Equivalent definition for mutual information

Claim 3

Let $(X, Y) \sim p$, then $I(X; Y) = D(p \| p_X p_Y)$.

► Proof:

$$\begin{aligned} D(p \| p_X p_Y) &= \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p_X(x) p_Y(y)} \\ &= \sum_{x,y} p(x, y) \log \frac{p_{X|Y}(x|y)}{p_X(x)} \\ &= - \sum_{x,y} p(x, y) \log p_X(x) + \sum_{x,y} p(x, y) \log p_{X|Y}(x|y) \\ &= H(X) + \sum_y p_Y(y) \sum_x p_{X|Y}(x|y) \log p_{X|Y}(x|y) \\ &= H(X) - H(X|Y) = I(X; Y). \square \end{aligned}$$

► We will later relate the above two claims.

Section 4

Relation to Data Compression

Wrong code

Theorem 4

Let p and q be distributions over $[m]$, and let C be code with

$\ell(i) = |C(i)| = \left\lceil \log \frac{1}{q_i} \right\rceil$. Then

$$H(p) + D(p\|q) \leq \mathbb{E}_{i \leftarrow p} [\ell(i)] \leq H(p) + D(p\|q) + 1$$

- ▶ Recall that $H(q) \leq \mathbb{E}_{i \leftarrow q} [\ell(i)] \leq H(q) + 1$.
- ▶ Proof of upperbound (upperbound is proved similarly)

$$\begin{aligned} \mathbb{E}_{i \leftarrow p} [\ell(i)] &= \sum_i p_i \left\lceil \log \frac{1}{q_i} \right\rceil < \sum_i p_i (\log \frac{1}{q_i} + 1) \\ &= 1 + \sum_i p_i (\log \frac{p_i}{q_i} \frac{1}{p_i}) = 1 + \sum_i p_i (\log \frac{p_i}{q_i}) + \sum_i p_i (\log \frac{1}{p_i}) \\ &= 1 + D(p\|q) + H(p) \end{aligned}$$

- ▶ Can there be a (close) to optimal code for q that is better for p ? HW

Section 5

Conditional Relative Entropy

Conditional relative entropy

Definition 5

For two distributions p and q over $\mathcal{X} \times \mathcal{Y}$:

$$\begin{aligned} D(p_{\mathcal{Y}|\mathcal{X}} \| q_{\mathcal{Y}|\mathcal{X}}) &:= \sum_{x \in \mathcal{X}} p_{\mathcal{X}}(x) \cdot \sum_{y \in \mathcal{Y}} p_{\mathcal{Y}|\mathcal{X}}(y|x) \log \frac{p_{\mathcal{Y}|\mathcal{X}}(y|x)}{q_{\mathcal{Y}|\mathcal{X}}(y|x)} \\ &= \mathbb{E}_{(X,Y) \sim p(X,Y)} \left[\log \frac{p_{\mathcal{Y}|\mathcal{X}}(Y|X)}{q_{\mathcal{Y}|\mathcal{X}}(Y|X)} \right] \end{aligned}$$

- Let $(X_p, Y_p) \sim p$ and $(X_q, Y_q) \sim q$, then

$$D(p_{\mathcal{Y}|\mathcal{X}} \| q_{\mathcal{Y}|\mathcal{X}}) = \mathbb{E}_{x \leftarrow X_p} [D(Y_q|_{X_p=x} \| Y_q|_{X_q=x})]$$

- Example: $p =$

$X \backslash Y$	0	1
0	$\frac{1}{8}$	$\frac{1}{8}$
1	$\frac{1}{4}$	$\frac{1}{2}$

 $q =$

$X \backslash Y$	0	1
0	$\frac{1}{8}$	$\frac{1}{4}$
1	$\frac{1}{2}$	$\frac{1}{8}$

$$\begin{aligned} D(p_{\mathcal{Y}|\mathcal{X}} \| q_{\mathcal{Y}|\mathcal{X}}) &= \frac{1}{4} \cdot D\left(\left(\frac{1}{2}, \frac{1}{2}\right) \parallel \left(\frac{1}{3}, \frac{2}{3}\right)\right) + \frac{3}{4} \cdot D\left(\left(\frac{1}{3}, \frac{2}{3}\right) \parallel \left(\frac{4}{5}, \frac{1}{5}\right)\right) \\ &= \dots \end{aligned}$$

Chain rule

Claim 6

For any two distributions p and q over $\mathcal{X} \times \mathcal{Y}$, it holds that
 $D(p\|q) = D(p_{\mathcal{X}}\|q_{\mathcal{X}}) + D(p_{\mathcal{Y}|\mathcal{X}}\|q_{\mathcal{Y}|\mathcal{X}})$

Proof:

$$\begin{aligned} D(p\|q) &= \sum_{(x,y) \in \mathcal{X} \times \mathcal{Y}} p(x,y) \log \frac{p(x,y)}{q(x,y)} \\ &= \sum_{(x,y) \in \mathcal{X} \times \mathcal{Y}} p(x,y) \log \frac{p_{\mathcal{X}}(x)p_{\mathcal{Y}|\mathcal{X}}(y|x)}{q_{\mathcal{X}}(x)q_{\mathcal{Y}|\mathcal{X}}(y|x)} \\ &= \sum_{(x,y) \in \mathcal{X} \times \mathcal{Y}} p(x,y) \log \frac{p_{\mathcal{X}}(x)}{q_{\mathcal{X}}(x)} + \sum_{(x,y) \in \mathcal{X} \times \mathcal{Y}} p(x,y) \log \frac{p_{\mathcal{Y}|\mathcal{X}}(y|x)}{q_{\mathcal{Y}|\mathcal{X}}(y|x)} \\ &= D(p_{\mathcal{X}}\|q_{\mathcal{X}}) + D(p_{\mathcal{Y}|\mathcal{X}}\|q_{\mathcal{Y}|\mathcal{X}}) \square \end{aligned}$$

Hence, for $(X, Y) \sim p$:

$$\begin{aligned} I(X, Y) &= D(p\|p_X p_Y) = D(p_X\|p_X) + \mathbb{E}_{x \leftarrow X} [D(p_{Y|X=x}\|p_Y)] \\ &= \mathbb{E}_{x \leftarrow X} [D(p_{Y|X=x}, p_Y)] \dots \end{aligned}$$

Section 6

Data-processing inequality

Data-processing inequality

Claim 7

For any rv's X and Y and function f , it holds that $D(f(X)\|f(Y)) \leq D(X\|Y)$.

- ▶ Analogues to $H(X) \geq H(f(X))$
- ▶ Proof:
- ▶ $D(X, f(X)\|Y, f(Y)) = D(X\|Y)$
- ▶ $D(X, f(X)\|Y, f(Y)) = D(f(X)\|f(Y)) + \mathbb{E}_{z \leftarrow f(X)} [D(X|_{f(X)=z}\|Y|_{f(X)=z})] \geq D(f(X)\|f(Y))$
- ▶ Hence, $D(f(X)\|f(Y)) \leq D(X\|Y)$.

Section 7

Relation to Statistical Distance

Relation to statistical distance

- ▶ $D(p||q)$ is used many time to measure the distance from p to q
- ▶ It is **not** a distance in the mathematical sense: $D(p||q) \neq D(q||p)$ and no triangle inequality
- ▶ However,

Theorem 8

$$SD(p, q) \leq \sqrt{\frac{\ln 2}{2} \cdot D(p||q)}$$

- ▶ Corollary: For rv X over $[m]$ with $H(X) \geq \log m - \varepsilon$, it holds that
$$SD(X, \sim [m]) \leq \sqrt{\frac{\ln 2}{2} \cdot (\log m - H(X))} = \sqrt{\frac{\ln 2}{2} \cdot \varepsilon}$$
- ▶ Other direction is incorrect: $SD(p, q)$ might be small but $D(p||q) = \infty$
- ▶ Does $SD(p, \sim [m])$ being small imply $D(p|| \sim [m]) = \log m - H(p)$ is small?

HW

Proving Thm 8, boolean case

► Let $p = (\alpha, 1 - \alpha)$ and $q = (\beta, 1 - \beta)$ and assume $\alpha \geq \beta$

► $SD(p, q) = \alpha - \beta$

► We will show that

$$D(p\|q) = \alpha \log \frac{\alpha}{\beta} + (1 - \alpha) \log \frac{1 - \alpha}{1 - \beta} \geq \frac{4}{2 \ln 2} (\alpha - \beta)^2 = \frac{2}{\ln 2} SD(p, q)^2$$

► Let $g(x, y) = x \log \frac{x}{y} + (1 - x) \log \frac{1 - x}{1 - y} - \frac{4}{2 \ln 2} (x - y)^2$

$$\begin{aligned} \frac{\partial g(x, y)}{\partial y} &= -\frac{x}{y \ln 2} + \frac{1 - x}{(1 - y) \ln 2} - \frac{4}{2 \ln 2} 2(y - x) \\ &= \frac{y - x}{y(1 - y) \ln 2} - \frac{4}{\ln 2} (y - x) \end{aligned}$$

► Since $y(1 - y) \leq \frac{1}{4}$, $\frac{\partial g(x, y)}{\partial y} \leq 0$ for $y < x$.

► Since $g(x, x) = 0$, $g(x, y) \geq 0$ for $y < x$. \square

Proving Thm 8, general case

- ▶ Let $\mathcal{U} = \text{Supp}(p) \cup \text{Supp}(q)$
- ▶ Let $\mathcal{S} = \{u \in \mathcal{U} : p(u) > q(u)\}$
- ▶ $\text{SD}(p, q) = \Pr_p[\mathcal{S}] - \Pr_q[\mathcal{S}]$ (by homework)
- ▶ Let $P \sim p$, and let the indicator \hat{P} be 1 iff $P \in \mathcal{S}$.
- ▶ Let $Q \sim q$, and let the indicator \hat{Q} be 1 iff $Q \in \mathcal{S}$.
- ▶ $\text{SD}(\hat{P}, \hat{Q}) = \Pr[P \in \mathcal{S}] - \Pr[Q \in \mathcal{S}] = \text{SD}(p, q)$

$$\begin{aligned} D(p \| q) &\geq D(\hat{P} \| \hat{Q}) && \text{(data-processing inequality)} \\ &\geq \frac{2}{\ln 2} \cdot \text{SD}(\hat{P}, \hat{Q})^2 && \text{(the Boolean case)} \\ &= \frac{2}{\ln 2} \cdot \text{SD}(p, q)^2. \quad \square \end{aligned}$$

Section 8

Conditioned Distributions

Main theorem

Theorem 9

Let X_1, \dots, X_k be iid over \mathcal{U} , and let $Y = (Y_1, \dots, Y_k)$ be rv over \mathcal{U}^k . Then $\sum_{j=1}^k D(Y_j \| X_j) \leq D(Y \| (X_1, \dots, X_k))$.

For rv Z , let $Z(z) = \Pr[Z = z]$.

We prove for $k = 2$, general case follows similar lines. Let $X = (X_1, X_2)$

$$\begin{aligned} D(Y \| X) &= \sum_{\mathbf{y} \in \mathcal{U}^2} Y(\mathbf{y}) \log \frac{Y(\mathbf{y})}{X(\mathbf{y})} = \sum_{\mathbf{y}=(y_1, y_2)} Y(\mathbf{y}) \log \frac{Y_1(y_1)}{X_1(y_1)} \frac{Y_2(y_2)}{X_2(y_2)} \frac{Y(\mathbf{y})}{Y_1(y_1) Y_2(y_2)} \\ &= \sum_{\mathbf{y}=(y_1, y_2)} Y(\mathbf{y}) \log \frac{Y_1(y_1)}{X_1(y_1)} + \sum_{\mathbf{y}=(y_1, y_2)} Y(\mathbf{y}) \log \frac{Y_2(y_2)}{X_2(y_1)} \\ &\quad + \sum_{\mathbf{y}=(y_1, y_2)} Y(\mathbf{y}) \log \frac{Y(\mathbf{y})}{Y_1(y_1) Y_2(y_2)} \\ &= D(Y_1 \| X_1) + D(Y_2 \| X_2) + I(Y_1; Y_2) \geq D(Y_1 \| X_1) + D(Y_2 \| X_2) \end{aligned}$$

Conditioning distributions, relative entropy case

Theorem 10

Let X_1, \dots, X_k be iid over \mathcal{X} , let $X = (X_1, \dots, X_k)$ and let W be an event (i.e., Boolean rv). Then $\sum_{j=1}^k D((X_j|_W) \| X_j) \leq D((X|_W) \| X) \leq \log \frac{1}{\Pr[W]}$.

$$\sum_{j=1}^k D((X_j|_W) \| X_j) \leq D((X|_W) \| X) \quad (\text{Thm 9})$$

$$\begin{aligned} &= \sum_{\mathbf{x} \in \mathcal{X}^k} (X|_W)(\mathbf{x}) \log \frac{(X|_W)(\mathbf{x})}{X(\mathbf{x})} \\ &= \sum_{\mathbf{x} \in \mathcal{X}^k} (X|_W)(\mathbf{x}) \log \frac{\Pr[W|X = \mathbf{x}]}{\Pr[W]} \quad (\text{Bayes}) \\ &= \log \frac{1}{\Pr[W]} + \sum_{\mathbf{x} \in \mathcal{X}^k} (X|_W)(\mathbf{x}) \log \Pr[W|X = \mathbf{x}] \\ &\leq \log \frac{1}{\Pr[W]} \end{aligned}$$

Conditioning distributions, statistical distance case

Theorem 11

Let X_1, \dots, X_k be iid over \mathcal{X} and let W be an event. Then

$$\sum_{j=1}^k \text{SD}((X_j|_W), X_j)^2 \leq \log \frac{1}{\Pr[W]}.$$

Proof: follows by Thm 8, and Thm 9. \square

Using $(\sum_{j=1}^k a_j)^2 \leq k \cdot \sum_{j=1}^k a_j^2$, it follows that

Corollary 12

$$\sum_{j=1}^k \text{SD}((X_j|_W), X_j) \leq \sqrt{k \log(\frac{1}{\Pr[W]})}, \text{ and}$$
$$\mathbb{E}_{j \leftarrow k} \text{SD}((X_j|_W), X_j) \leq \sqrt{\frac{1}{k} \log(\frac{1}{\Pr[W]})}$$

Extraction

Numerical example

- ▶ Let $X = (X_1, \dots, X_k) \leftarrow \{0, 1\}^{40}$ and let $f: \{0, 1\}^{40} \mapsto 0$ be such that $\Pr[f(X) = 0] = 2^{-10}$.
- ▶ $E_{j \leftarrow [40]} \text{SD}((X_j|_{f(X)=0}), \sim \{0, 1\}) \leq \sqrt{\frac{1}{40} \cdot 10} = \frac{1}{2}$
- ▶ Typical bits are not too biased, even when conditioning on a very unlikely event.

Extension

Theorem 13

Let $X = (X_1, \dots, X_k)$, T and V be rv's over \mathcal{X}^k , \mathcal{T} and \mathcal{V} respectively. Let W be an event and assume that the X_i 's are iid conditioned on T . Then

$$\sum_{j=1}^k D((TVX_j)|_w || (TV)|_w X'_j(T)) \leq \log \frac{1}{\Pr[W]} + \log |\text{Supp}(V|_w)|,$$

where $X'_j(t)$ is distributed according to $X_j|_{T=t}$.

Interpretation.

Proving Thm 13

Let $X = (X_1, \dots, X_k)$, T and V be rv's over \mathcal{X}^k , \mathcal{T} and \mathcal{V} respectively, such that X_i 's are iid conditioned on T . Let W be an event and let $X_j'(t)$ be distributed according to the distribution of $X_j|_{T=t}$.

$$\begin{aligned} & \sum_{j=1}^k D((TVX_j)|_W || (TV)|_W X_j'(T)) \\ &= \mathbb{E}_{(t,v) \leftarrow (TV)|_W} \left[\sum_{j=1}^k D(X_j|_{W, V=v, T=t} || (X_j|_{T=t})) \right] \end{aligned} \quad \text{(chain rule)}$$

$$\begin{aligned} &= \mathbb{E}_{(t,v) \leftarrow (TV)|_W} \left[\sum_{j=1}^k D(\underbrace{(X_j|_{W, V=v})}_{W'}|_{T=t} || (X_j|_{T=t})) \right] \\ &\leq \mathbb{E}_{(t,v) \leftarrow (TV)|_W} \left[\log \frac{1}{\Pr[W \wedge V=v | T=t]} \right] \end{aligned} \quad \text{(Thm 10)}$$

$$\leq \log \mathbb{E}_{(t,v) \leftarrow (TV)|_W} \frac{1}{\Pr[W \wedge V=v | T=t]} \quad \text{(Jensen's inequality)}$$

$$= \log \sum_{(t,v) \in \text{Supp}((TV)|_W)} \frac{\Pr[T=t]}{\Pr[W]} \leq \log \frac{|\text{Supp}(V|_W)|}{\Pr[W]}.$$

□